

Vishal Singh Yadav

LinkedIn: [linkedin.com/in/tokudai/](https://www.linkedin.com/in/tokudai/)

Github: github.com/tokudai-vs

Address: NCR / Hyderabad, India

Email: vishal.singh@alumni.iith.ac.in

Mobile: +91-9555756235

SUMMARY

Data Scientist (Associate DS 3) with 5+ years of combined experience — 3 years in applied research and 2+ years driving scalable industry solutions in NLP, GenAI, and risk modeling. Experienced in building and deploying end-to-end ML pipelines involving LLMs (LLAMA, RAG, QLoRA), time-series forecasting, and NL-to-SQL systems, delivering \$1.5M+ business impact and enabling 15x growth in targeted outreach. Skilled in deploying robust solutions on AWS, Azure, and hybrid cloud infrastructure, primarily in healthcare and insurance domains.

EXPERIENCE

- **Carelon Global Solutions (*Elevance Health*)** Hyderabad, India
Associate Data Scientist III *Jul 2023 – Present*
 - Designed and shipped forecasting modules using anomaly detection and GenAI for the Cost of Care platform, converting large-scale claims data (60M+ members) into 3-month state-level predictions to support proactive healthcare interventions.
 - Worked on NLP-based feature engineering and optimisation using financial, clinical, and demographic data to reduce underwriting turnaround by 20% and improve real-time rating accuracy — contributing to projected \$1.5M in cost savings by 2025.
 - Developed probabilistic models to flag high-cost claimants and special medical conditions, enabling \$300K in annual savings through better cost visibility and risk stratification. Improved condition identification coverage from 48% to 81%, driving a 15x increase in outreach scale.
 - Built a NL-to-SQL pipeline using fine-tuned LLMs, semantic search, and visual query flows — delivered 95% accuracy and improved performance over prior methods by 31%.
 - Co-developed a clinical QA assistant with LLAMA2 and LangChain using a RAG pipeline and QLoRA-tuned base model — enabled secure, domain-specific medical question answering at scale.
 - Automated ticket triage workflows with Word2Vec+BERT embeddings and XGBoost — eliminated repetitive manual classification, saving 8–10 FTEs worth of effort across support teams.
 - Contributed to a 40% improvement in underwriting decision throughput by developing a clinical rules engine, integrating ML-based risk scoring, and LLM-assisted automation across workflows — delivering \$100K+ in operational impact.
 - Supported deployment of ML models across AWS, Azure, and on-prem environments using Docker-based pipelines — enabled stable rollout of real-time APIs and dashboard integrations.
- **Qulabs Software India Pvt Ltd** Hyderabad, India
Machine Learning Engineer *Mar 2023 – Jul 2023*
 - Built a RAG-style semantic search QA chatbot using BERT embeddings and vector search over 100K+ files — enabled fast, context-relevant answers via semantic similarity, without using LLMs.
 - Mentored 3 interns and 3 junior developers on ML workflows, code quality, and model deployment — improved maintainability and established reusable standards across the team.
- **Krama Lab, IIT Hyderabad (*Machine Learning Research*)** Hyderabad, India
Research Assistant *Jan 2020 – Dec 2022*
 - **Industry Collaboration (GreatFour Systems):** Conducted independent research on point cloud segmentation, graph neural networks (GNNs), and few-shot learning under faculty guidance — developed core model architectures for structured reasoning and sparse data classification in industrial settings.
 - **Academic Research:** Designed and evaluated experimental GNN-based models for knowledge graph completion, link prediction, and representation learning under low-resource constraints.

EDUCATION

- **Indian Institute of Technology Hyderabad** Hyderabad, India
M.Tech (by Research) in Computer Science and Engineering GPA: 8.38 / 10 *Jan 2020 – Dec 2022*
- **Guru Gobind Singh Indraprastha University** New Delhi, India
B.Tech in Computer Science and Engineering Percentage: 71.2% *Jul 2014 – May 2018*

SKILLS

- **Generative AI, NLP & Machine Learning:** LLAMA, QLoRA, Fine-tuning, Prompt Engineering, RAG, AI Agents, MCP (Model Context Protocol), Transformers, BERT, spaCy, Semantic Search, Graph Neural Networks, Knowledge Graphs, Few-Shot Learning, Anomaly Detection, PyTorch, TensorFlow, Scikit-learn, XGBoost
- **MLOps & Cloud:** AWS (EC2, S3, SageMaker), Docker, CI/CD, MLflow, Weights & Biases, FastAPI, Model Serving, ONNX, A/B Testing, Performance Tuning, Streamlit, Gradio
- **Data Engineering & Tools:** PySpark, Snowflake, Distributed Computing, Feature Stores, ETL Pipelines, Time-Series Analysis, Data Augmentation, FAISS
- **Programming Languages:** Python, SQL, C++ (STL, Boost), Bash

PROJECTS

- **CodeChakra: AI-Powered PR Review Assistant (OSS, WIP):** Building a secure, end-to-end AI assistant to automate pull request reviews by simulating full-stack engineering audits. Triggered on GitHub PRs, it performs static analysis, security checks, latency profiling (I/O, DB, network ops), and suggests performance reliability improvements via LLMs. Designed for high-trust environments with encrypted I/O, modular agent orchestration, and planned end to end encryption support. Includes CLI + CI/CD integration hooks for DevOps workflows, and plug-and-play architecture for future ML-powered explainability, anomaly detection, and repo-level performance insights. Open-source and extensible — PoC in progress with Docker-based local deployment and multi-agent coordination in pipeline.

ACHIEVEMENTS

- **Go Above Award (x2) – Carelon Global Solutions:** Twice awarded for high-impact execution under critical timelines — launched a member insights dashboard within 24 hours and led the expedited delivery of the ARROW project during peak business pressure.
- **Rapid Product Deployment – Carelon Global Solutions:** Personally awarded by the Director of Data Science for leading the fast-tracked build and deployment of a key analytics product — accelerated insight delivery and enabled smooth production rollout.
- **Mentor – AI and Advanced Technologies:** Delivered technical mentorship to working professionals through the IIT Hyderabad & NSE Talentsprint program — trained cohorts in core ML workflows and building production-grade GenAI systems.

CERTIFICATIONS

- **Advanced Certification in Generative AI – Carelon / Prizmato:** Covered production deployment of LLMs, QLoRA fine-tuning, RAG system design, HIPAA-compliant prompt engineering, and optimization techniques for GenAI workflows.
- **NVIDIA Deep Learning Institute: NLP & Deep Learning:** Fundamentals of Deep Learning, Transformer-Based NLP Applications
- **AI for Ops:** AI for Anomaly Detection, Accelerated Data Engineering Pipelines
- **Accelerated Computing:** Fundamentals of Accelerated Computing and Data Science

PUBLICATIONS

- **Context-Aware Question Routing in Community Question Answering Sites:** *Vishal Singh Yadav, Manish Singh May 2023.* Proposed a deep learning framework for intelligent question routing using contextual embeddings and semantic similarity scoring to improve query-to-expert mapping.